

AgriVision: Automated Multi-Class Fruit Classification and Quality Control System using Deep Transfer Learning

1. Project Overview

In the agricultural industry and food supply chains, automated quality control is essential to reduce waste and ensure safety. Your task is to build a computer vision system capable of classifying fruits.

This project focuses on modernizing agricultural supply chains by developing a robust Computer Vision system capable of identifying distinct fruit varieties. By leveraging **Transfer Learning** (using architectures like MobileNetV2 or ResNet), you will build a scalable, high-accuracy model to automate manual sorting processes, reducing waste and increasing efficiency in food processing lines.

2. Learning Objectives

By completing this project, you will demonstrate the ability to:

- Load and preprocess image datasets for deep learning.
- Implement data augmentation strategies to improve model generalization.
- Apply Transfer Learning using pre-trained architectures (e.g., VGG16, ResNet, MobileNet).
- Fine-tune models to adapt to specific domain tasks.
- Evaluate model performance using appropriate metrics and visualizations.

3. Dataset Details

You must use the following dataset for this project. No external datasets are allowed unless explicitly permitted for testing purposes.

- Source: Roboflow - <https://universe.roboflow.com/nunnapas/fruit-rzujm>
- Classes: The dataset covers 10 types of fruits: Apple, Avocado, Banana, Cherry, Kiwi, Mango, Orange, Pineapple, Strawberries, and Watermelon.
- Structure:
 - test/ (Contains images for evaluating the model)
 - train/ (Contains images for training the model)
 - valid/ (Contains validation data)

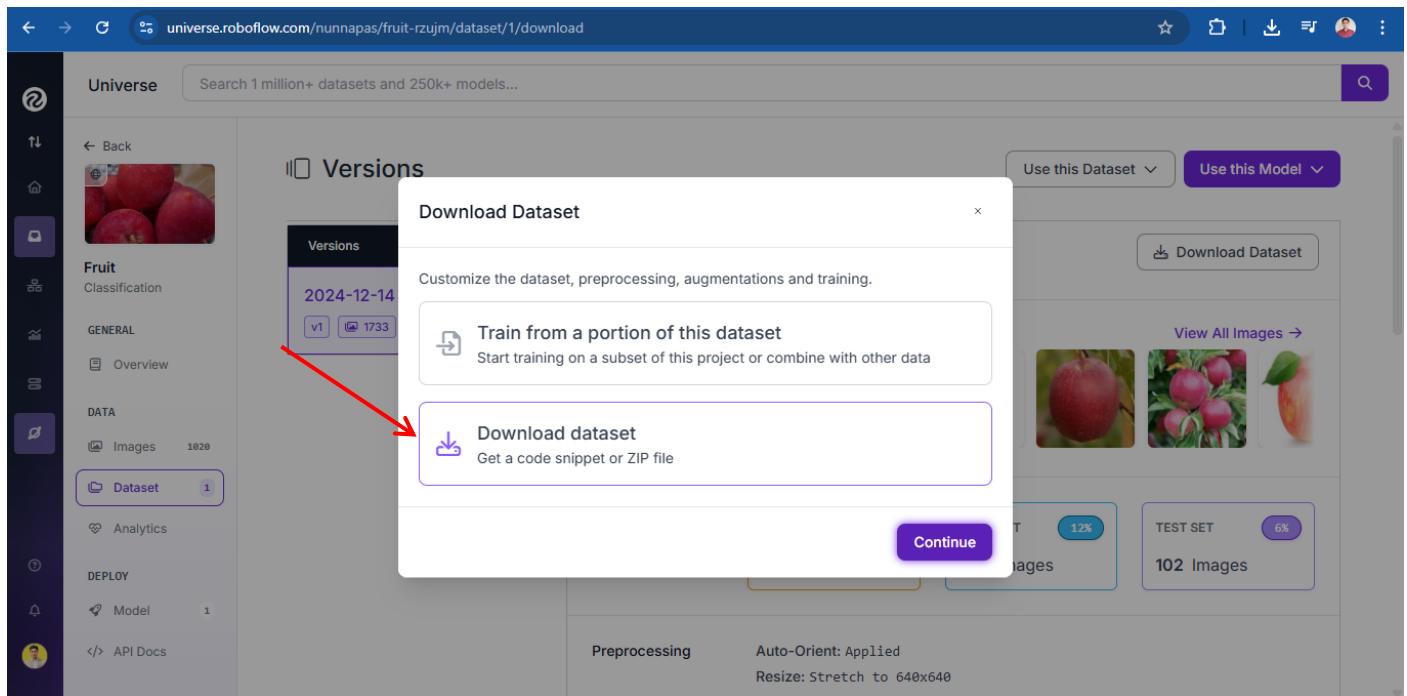
Here is how you can download the data: Click [Here](#) > Under Data section select Dataset

The screenshot shows the Roboflow Universe interface. On the left, there's a sidebar with navigation icons and a list of sections: Back, Fruit (Classification), GENERAL (Overview, Dataset, Images, Analytics), DATA (Images, 1820, Dataset 1, highlighted with a red box), DEPLOY (Model 1), and API Docs. The main content area displays the 'Fruit Computer Vision Model' by Nunnapas, updated a year ago. It features several fruit images (apple, watermelon, pineapple, strawberries) and a summary card for the model. The 'Dataset' section in the sidebar has a red arrow pointing to it.

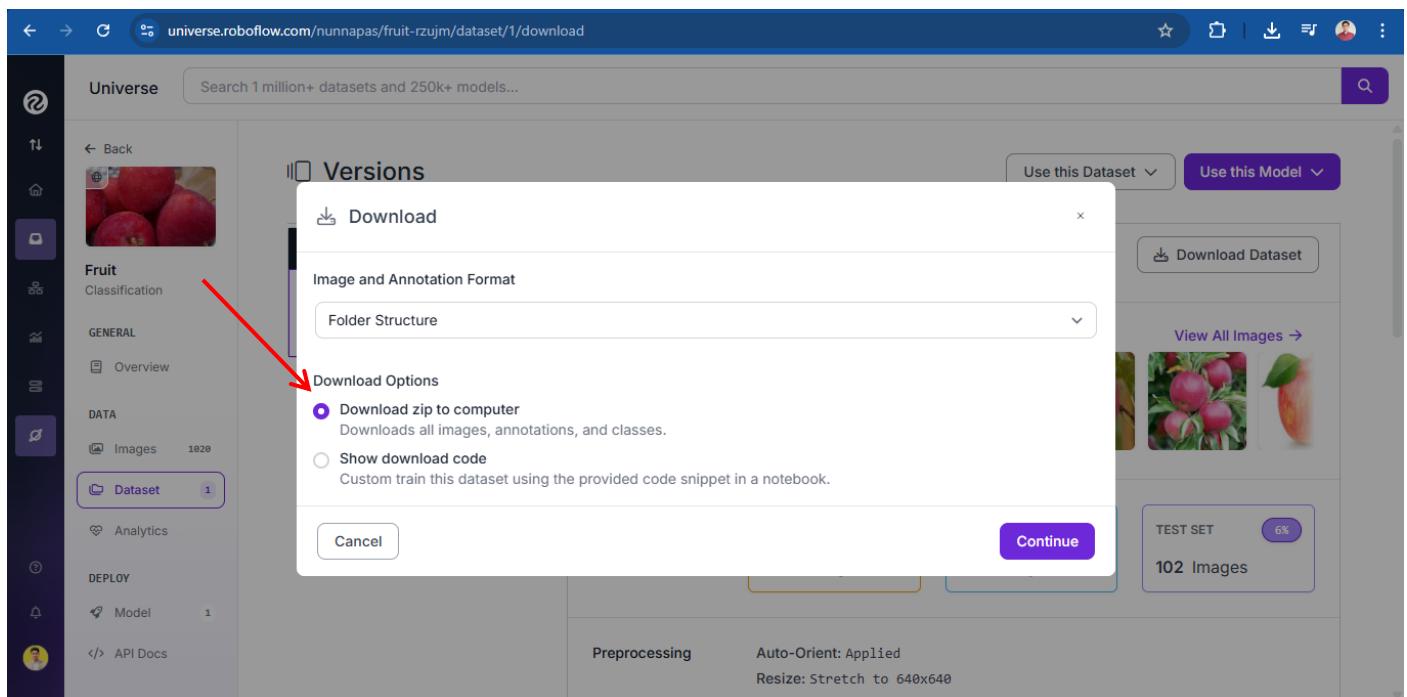
Click on Download Dataset Button

The screenshot shows the 'Versions' page for the dataset. The sidebar remains the same. The main content area shows a list of versions: v1 (selected, highlighted with a red box) and v2 (2024-12-14 11:20am). Below this, there are sections for '1733 Total Images' (with preview images) and 'Dataset Split' (Train Set: 1426 Images, 82%; Valid Set: 205 Images, 12%; Test Set: 102 Images, 6%). A red arrow points to the 'Download Dataset' button in the top right corner of the main content area.

Select Download Dataset and Click Continue



Select Download zip to computer and Continue. Hurrah! Check the downloads directory on your machine.



- Upload the ZIP file to your Google Drive. Mount your Drive in the notebook (from google.colab import drive) and unzip the dataset programmatically using Python (e.g., !unzip).
- If using Local Jupyter: Extract the ZIP file into your project directory.

4. Project Tasks & Methodology

You are expected to follow the machine learning pipeline described below.

Phase 1: Data Exploration & Preprocessing (15 Points)

1. Exploratory Data Analysis (EDA):

- Identify Classes: Use Python's os library to list the sub-folders in the train directory. Print the list of class names.
- Visualize Samples: Display one sample image for every class found.
- Distribution Plot: Plot a bar chart showing the number of images per class in the train set to check for imbalance.

2. Preprocessing:

- Resize images to the input shape required by your chosen model (e.g., 224x224).
- Normalize pixel values (e.g., rescale to 0-1 range).

3. Data Augmentation:

- Implement data augmentation on the **training set only** to prevent overfitting.
- *Suggested techniques:* Rotation, Zoom, Horizontal Flip, Width/Height Shift.

Phase 2: Model Development - Transfer Learning (40 Points)

You must implement a robust Transfer Learning solution.

1. Select a Pre-trained Model:

- Choose a standard architecture (e.g., VGG16, ResNet50, InceptionV3, or MobileNetV2) pre-trained on ImageNet.

2. Feature Extraction:

- Load the pre-trained model **without the top** (fully connected) classification layers (often called include_top=false).
- Freeze the weights of the base layers so they are not updated during the initial training.

3. Custom Classification Head:

- Add your own dense layers with ReLU activation on top of the frozen base.
- Include Dropout layers to reduce overfitting.
- **Output Layer:** Must use Softmax activation with N units (N is number of distinct classes).

4. Training:

- Compile the model using Adam optimizer and CategoricalCrossentropy loss function.
- Train for enough epochs (minimum 5 or 10).
- Implement **Early Stopping** to prevent unnecessary computation if the validation loss stops improving.

Phase 3: Fine-Tuning (Optional)

- Once the head is trained, you may unfreeze some of the top layers of the pre-trained base model.
- Re-train the model with a very low learning rate (e.g., 1e-5) to adapt specific features to the fruit dataset.

Phase 4: Evaluation & Analysis (30 Points)

1. Quantitative Metrics:

- Report the final Accuracy on the Test set.
- Generate a Classification Report (Precision, Recall, F1-Score for each class).

2. Visualizations:

- Training Curves:** Plot Training vs. Validation Accuracy and Loss over epochs.
- Confusion Matrix:** Plot a heatmap of the confusion matrix

3. Error Analysis:

- Display 5 images from the test set where the model made a **wrong prediction**.
- Provide a brief comment on *why* the failure might have happened (e.g., "The model confused a Red Apple with a Tomato due to similar shape and color").

Phase 5: Deployment / Inference (15 Points)

- Create a simple function predict_fruit(image_path) that takes a local image path as input and prints:
 - The predicted class.
 - Returns the **Predicted Class** and the **Confidence Score** (probability).

5. Deliverables

You must submit the following:

- Source Code (.ipynb):** A clean, runnable Google Colab notebook.
 - Code must be well-commented.
 - Output cells (graphs, logs) must be visible (do not clear outputs).
- Project Report (PDF):** A 2-3-page summary containing:
 - Introduction & Objective.
 - Methodology (Which model did you choose and why?).
 - Results (Screenshots of graphs, confusion matrix, and metrics).
 - Conclusion (What worked, what didn't?).